

Setup of mong and elasticsearch

Clément Levallois

2017-04-10

Table of Contents

System	1
Why MongoDB and Elasticsearch?	1
Installing MongoDB	1
Install Elasticsearch	3
Config	4
Install the mongo to elasticsearch connection	4
elastic2-doc-manager	4
run Mongo as a replicaset:	5
Install kibana	5
Install X-pack	6
Connect Mongo to ElasticSearch:	6
the end	6

last modified: 2017-04-11

System

- I use Debian, version 8.7 ([why?](#))
- Vi is used as a text editor in the following
- **MongoDB 3.4 and Elasticsearch 5.x**

Why MongoDB and Elasticsearch?

- MongoDB is a database which stores data without the need for a pre-established model ("strict description") of this data. In practice: I can save something into MongoDB without spending time creating tables and stuff. Just save a JSON doc, that's it.
- MongoDB alone is great, but I will store gigabytes of data, with several text fields and some simple graph logic as well. Elasticsearch is known for managing well the indexes and queries related to these data types.
- A blog post which details how Elasticsearch helped on performances for Mongo: <http://blog.quarkslab.com/mongodb-vs-elasticsearch-the-quest-of-the-holy-performances.html>
- And Elasticsearch can then integrate with [Kibana](#), a way to visualize query results. Awesome!

Installing MongoDB

source: <https://docs.mongodb.com/manual/tutorial/install-mongodb-on-debian/>

```
sudo apt-key adv --keyserver hkp://keyserver.ubuntu.com:80 --recv  
0C49F3730359A14518585931BC711F9BA15703C6  
echo "deb http://repo.mongodb.org/apt/debian jessie/mongodb-org/3.4 main" | sudo tee  
/etc/apt/sources.list.d/mongodb-org-3.4.list
```

disable Transparent Huge Pages as per <https://docs.mongodb.com/manual/tutorial/transparent-huge-pages/>

```
sudo vi /etc/init.d/disable-transparent-hugepages
```

paste this in the text editor:

```
#!/bin/bash
### BEGIN INIT INFO
# Provides:          disable-transparent-hugepages
# Required-Start:    $local_fs
# Required-Stop:
# X-Start-Before:    mongod mongodbmms-automation-agent
# Default-Start:     2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Disable Linux transparent huge pages
# Description:       Disable Linux transparent huge pages, to improve
#                   database performance.
### END INIT INFO

case $1 in
  start)
    if [ -d /sys/kernel/mm/transparent_hugepage ]; then
      thp_path=/sys/kernel/mm/transparent_hugepage
    elif [ -d /sys/kernel/mm/redhat_transparent_hugepage ]; then
      thp_path=/sys/kernel/mm/redhat_transparent_hugepage
    else
      return 0
    fi

    echo 'never' > ${thp_path}/enabled
    echo 'never' > ${thp_path}/defrag

    re='^[0-1]+$'
    if [[ $(cat ${thp_path}/khugepaged/defrag) =~ $re ]]
    then
      # RHEL 7
      echo 0 > ${thp_path}/khugepaged/defrag
    else
      # RHEL 6
      echo 'no' > ${thp_path}/khugepaged/defrag
    fi

    unset re
    unset thp_path
    ;;
  esac
```

Make the file executable:

```
sudo chmod 755 /etc/init.d/disable-transparent-hugepages
```

Make the file to be ran on reboot:

```
sudo update-rc.d disable-transparent-hugepages defaults
```

Start Mongo:

-change for user mongod:

```
sudo -u mongod bash
```

Then only actually start mongo:

```
sudo service mongod start
```

Check that it runs fine:

```
sudo cat /var/log/mongodb/mongod.log
```

→ There should be a line "[initandlisten] waiting for connections on port <port>"

Install Elasticsearch

source: <https://www.elastic.co/guide/en/elasticsearch/reference/current/deb.html>

Download the public signing key:

```
wget -q0 - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -
```

(can take a while)

Then:

```
sudo apt-get install apt-transport-https
```

```
echo "deb https://artifacts.elastic.co/packages/5.x/apt stable main" | sudo tee -a  
/etc/apt/sources.list.d/elastic-5.x.list
```

```
sudo /bin/systemctl daemon-reload  
sudo /bin/systemctl enable elasticsearch.service
```

Config

```
switch to root user:
```

```
su -  
cd /etc/default/elasticsearch  
vi elasticsearch.yml
```

→ switch this param to true: bootstrap.memory_lock: true

You then need to make sure the JVM Heap size is no more than half the RAM. First fix a memory param:

```
sudo mkdir /etc/systemd/system/elasticsearch.service.d  
sudo vi elasticsearch.conf
```

Add these lines:

LimitMEMLOCK=infinity

Adjust resource limits:

```
sudo vi /etc/security/limits.conf
```

Add line:

elasticsearch - nofile 65536

Add a jvm parameter:

```
sudo /etc/elasticsearch/jvm.options
```

Add this line:

-Djava.io.tmpdir=/var/tmp

Install the mongo to elasticsearch connection

elastic2-doc-manager

This is a doc manager by mongodb labs.

Source: <https://github.com/mongodb-labs/elastic2-doc-manager>

```
sudo easy_install pip
sudo pip install 'elastic2-doc-manager[elastic5]'
sudo pip install 'mongo-connector[elastic5]'
```

run Mongo as a replicaset:

```
sudo service mongod stop
```

```
sudo vi /etc/mongodb.conf
```

Change dbPath to /data/db

Then:

```
sudo chown -R mongodb:mongodb /data/db
```

Then launch mongo as a replicaset:

```
sudo mongod --port 27017 --dbpath /data/db --replSet rs0 --fork --logpath
/var/log/mongodb.mongod.log
```

Install kibana

Kibana is the visualization engine for elastic.

```
sudo wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add
-
sudo apt-get install kibana
```

Configure Kibana to start automatically at boot:

```
sudo /bin/systemctl daemon-reload
sudo /bin/systemctl enable kibana.service
```

Start Kibana:

```
sudo systemctl start kibana.service
```

Install X-pack

<https://www.elastic.co/guide/en/x-pack/current/installing-xpack.html>

Connect Mongo to ElasticSearch:

Source: <https://blog.jixee.me/how-to-use-mongo-connector-with-elasticsearch/>

```
sudo mongo-connector -m localhost:27017 -t localhost:9200 -d elastic2_doc_manager -n database1.collection1,database1.collection2
```

You can check that the connection is made here, your Mongo collections should be listed on this page:

http://localhost:9200/_cat/indices?v

the end

Author of this tutorial: [Clement Levallois](#)

All resources on linux security: <https://seinecle.github.io/linux-security-tutorials/>